

Redshift is SUPER

A More Flexible Cloud Data Warehouse Design



SPECIAL THANKS TO ALL OUR AWESOME CAMP SPONSORS!



TruStage™



HEARTLAND
BUSINESS SYSTEMS



Intellibus



Unspecified
SOFTWARE CO



Architect.io



ionic

CAMUNDA



Google Cloud

COREBTS



GrapeCity®



Join Us
Next Year!

WI 24

THAT[®]
CONFERENCE

JULY 29TH - AUG. 1ST



My Story

Nick Heidke, Principal Data Architect at SimpleRose

Live in Appleton, WI

Bachelor's in Comp Sci
from University of
Wisconsin-Eau Claire 2009



Alright, so, catchy title Nick,
lots of buzzwords, well
done, but what do I
actually need to know so
this makes any sense?



Glad you asked. Couple of key points:

- The keyword in cloud data warehouse is warehouse meaning we're working with OLAP instead of OLTP
- Our specific technology, and several like it, is also capable of Massive Parallel Processing (MPP)
- While MPP databases are incredibly fast due to the amount of compute power available to them, our design must be cognizant of how the data is laid out



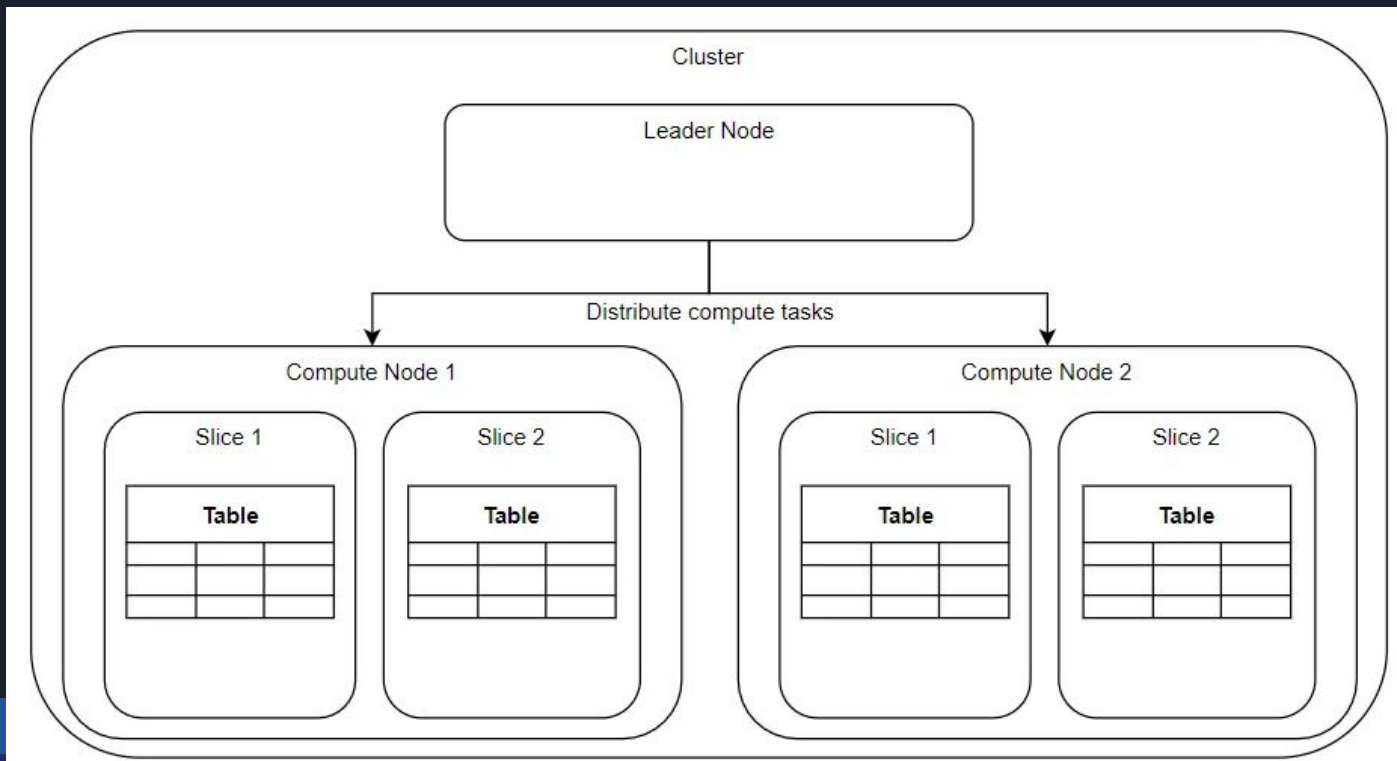


Key Features of MPP

- Parallel Processing - Data and queries are divided across nodes, each independently performing a chunk of the work
- Massive scale - On the order of petabytes
- Shared nothing architecture
- Horizontal instead of vertical scaling
- Distribution and sort keys - no traditional indexes



MPP Architecture



Schema Design with MPP In Mind

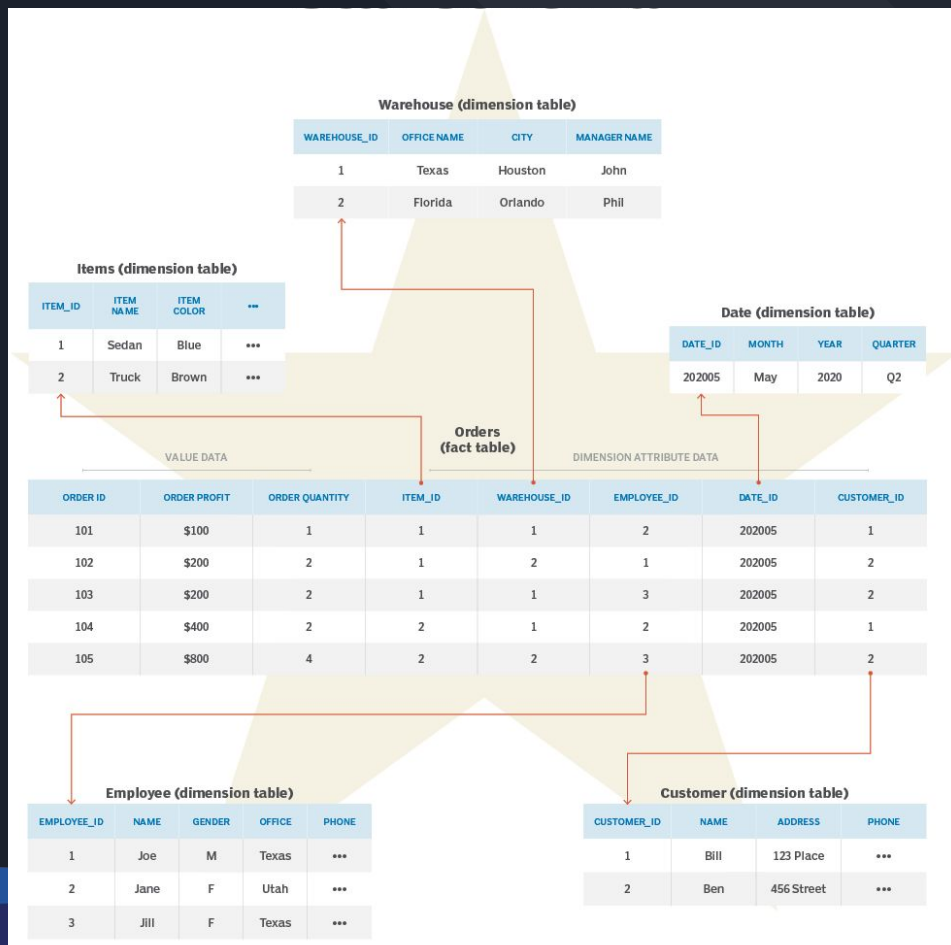
- Typically, a single table's data is distributed across all of the nodes
- Data on each node is sorted by one key (though it can include several columns)
- If required data is not on a given node, it must be copied from another
- Where does that lead us?



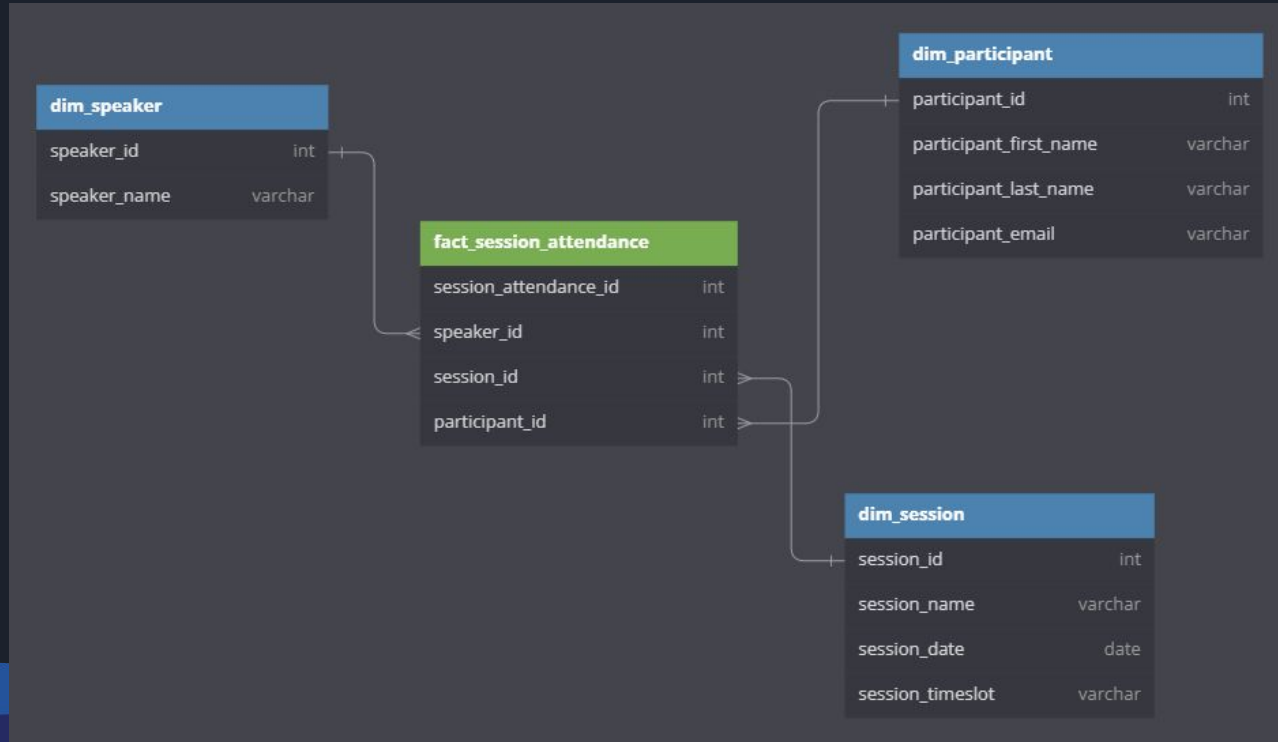
To the stars...



Specifically, the star schema

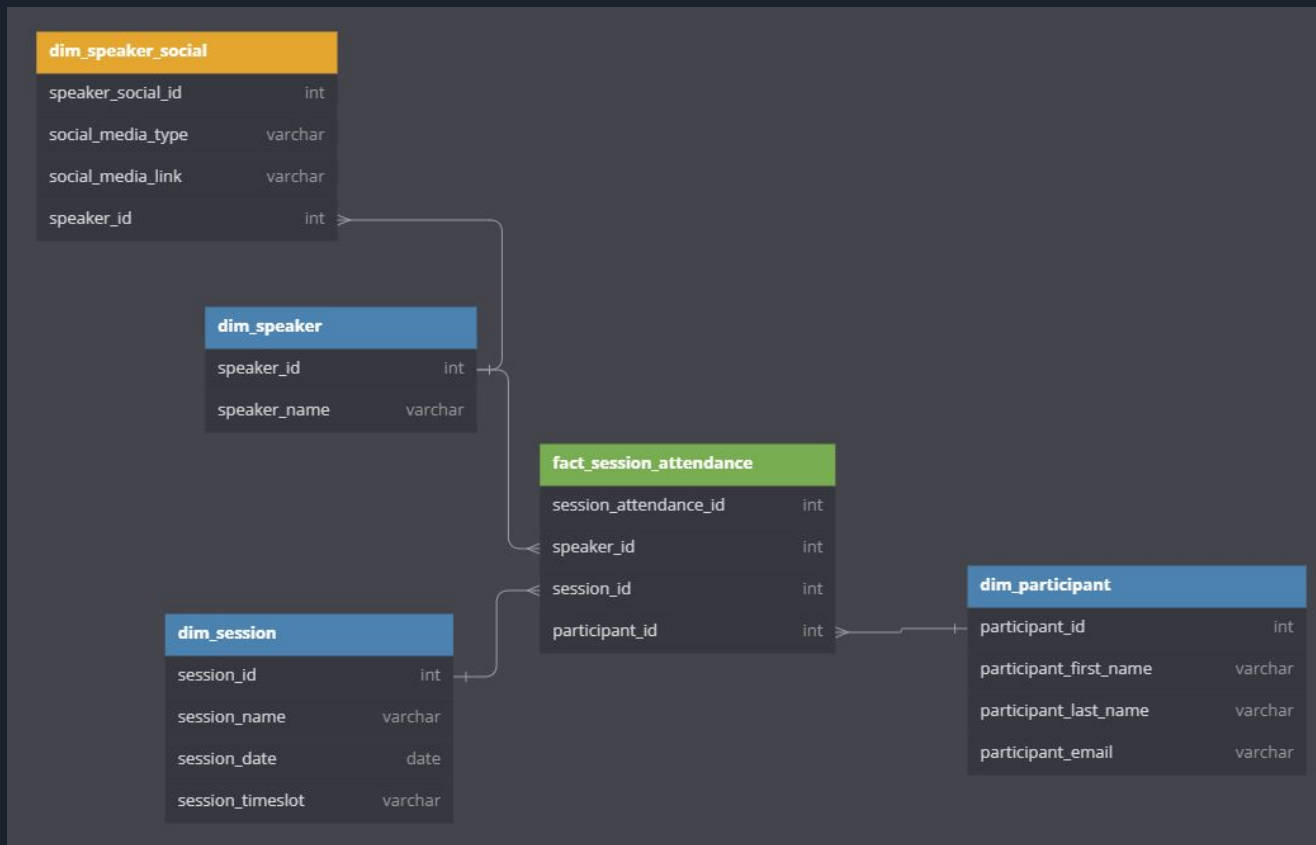


Our domain problem today



What happens when we
need to add social
media contacts for our
speakers?





We might start out by adding an additional table linked to our existing speaker dimension



Our first query requirement comes in

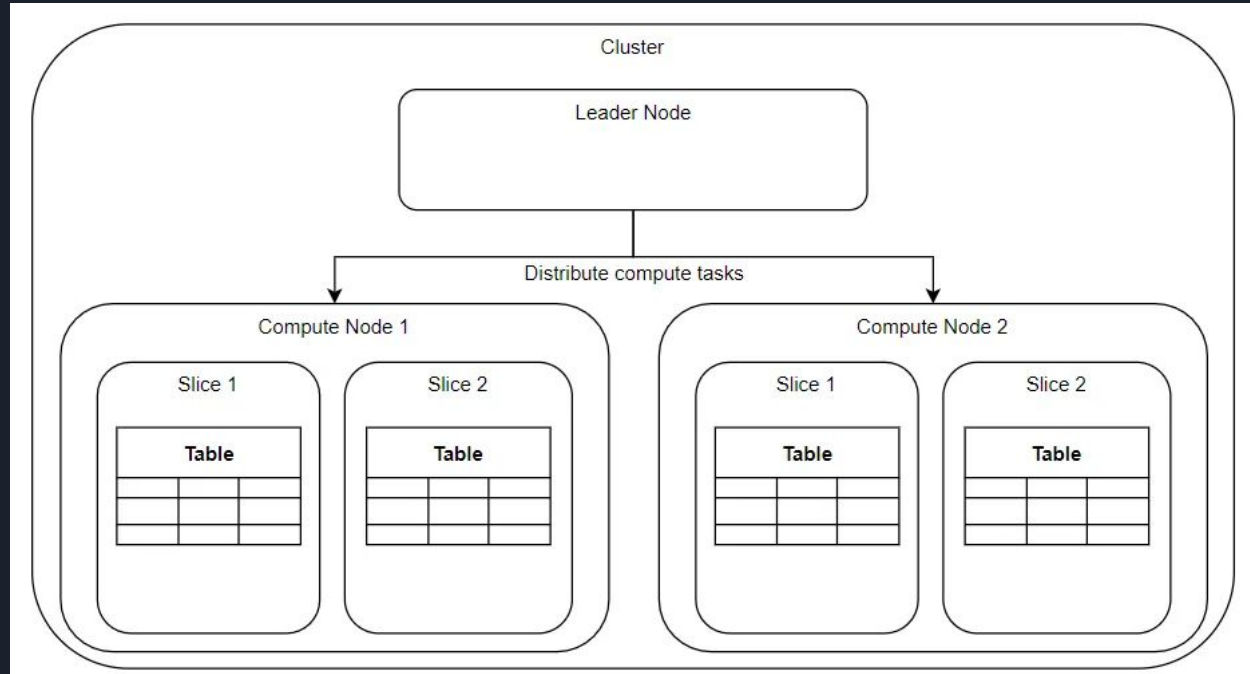
```
SELECT ds.speaker_name, dss.social_media_link, dss.social_media_type, d.session_name
FROM conf_star_schema_v1.fact_session_attendance fsa
INNER JOIN conf_star_schema_v1.dim_speaker ds
ON fsa.speaker_id = ds.speaker_id
INNER JOIN conf_star_schema_v1.dim_speaker_social dss
ON dss.speaker_id = ds.speaker_id
INNER JOIN conf_star_schema_v1.dim_session d
ON fsa.session_id = d.session_id
WHERE speaker_name = 'John Smith';
```



What's the potential issue?

Slices of data are
each on independent
nodes

Joining across the
nodes is expensive



Query Cost - Network Ops

Network operation	Distribution Behavior	Relative Cost
DS_DIST_NONE DS_DIST_ALL_NONE	No redistribution required	Low
DS_DIST_INNER DS_DIST_OUTER DS_DIST_BOTH	One or more of the tables is redistributed	Medium/High
DS_BCAST_INNER	A copy of the entire inner table is broadcast to all the compute nodes	High



How does our query plan look currently?

QUERY PLAN

```
XN Hash Join DS_BCAST_INNER (cost=640000.60..1240000.87 rows=2 width=89)
  Hash Cond: ("outer".speaker_id = "inner".speaker_id)
    -> XN Seq Scan on dim_speaker_social dss (cost=0.00..0.11 rows=11 width=51)
    -> XN Hash (cost=640000.60..640000.60 rows=1 width=50)
      -> XN Hash Join DS_BCAST_INNER (cost=280000.36..640000.60 rows=1 width=50)
        Hash Cond: ("outer".session_id = "inner".session_id)
          -> XN Seq Scan on dim_session d (cost=0.00..0.10 rows=10 width=31)
          -> XN Hash (cost=280000.36..280000.36 rows=1 width=27)
            -> XN Hash Join DS_BCAST_INNER (cost=0.13..280000.36 rows=1 width=27)
              Hash Cond: ("outer".speaker_id = "inner".speaker_id)
                -> XN Seq Scan on fact_session_attendance fsa (cost=0.00..0.10 rows=10 width=8)
                -> XN Hash (cost=0.12..0.12 rows=1 width=19)
                  -> XN Seq Scan on dim_speaker ds (cost=0.00..0.12 rows=1 width=19)
                  Filter: ((speaker_name)::text = 'John Smith'::text)
```



Option #1 - Distribution Styles

Distribution Style	Description
EVEN	Data is distributed across all slices evenly
KEY	Specified key is used to distribute data across slices; rows with the same distribution key are co-located within the same slice
ALL	Entire table is replicated to every slice in the cluster
AUTO	Allows Redshift to automatically determine optimal distribution based on table size and heuristics



What distribution style
is best for our case?



Table Design - Version 2

```
create table dim_speaker  
(  
    speaker_id    integer distkey,  
    speaker_name  varchar(256)  
) DISTSTYLE KEY;
```

```
create table dim_speaker_social  
(  
    speaker_social_id integer,  
    social_media_type  varchar(256),  
    social_media_link  varchar(256),  
    speaker_id         integer distkey  
) DISTSTYLE KEY;
```



Query Plan - Version 2

QUERY PLAN

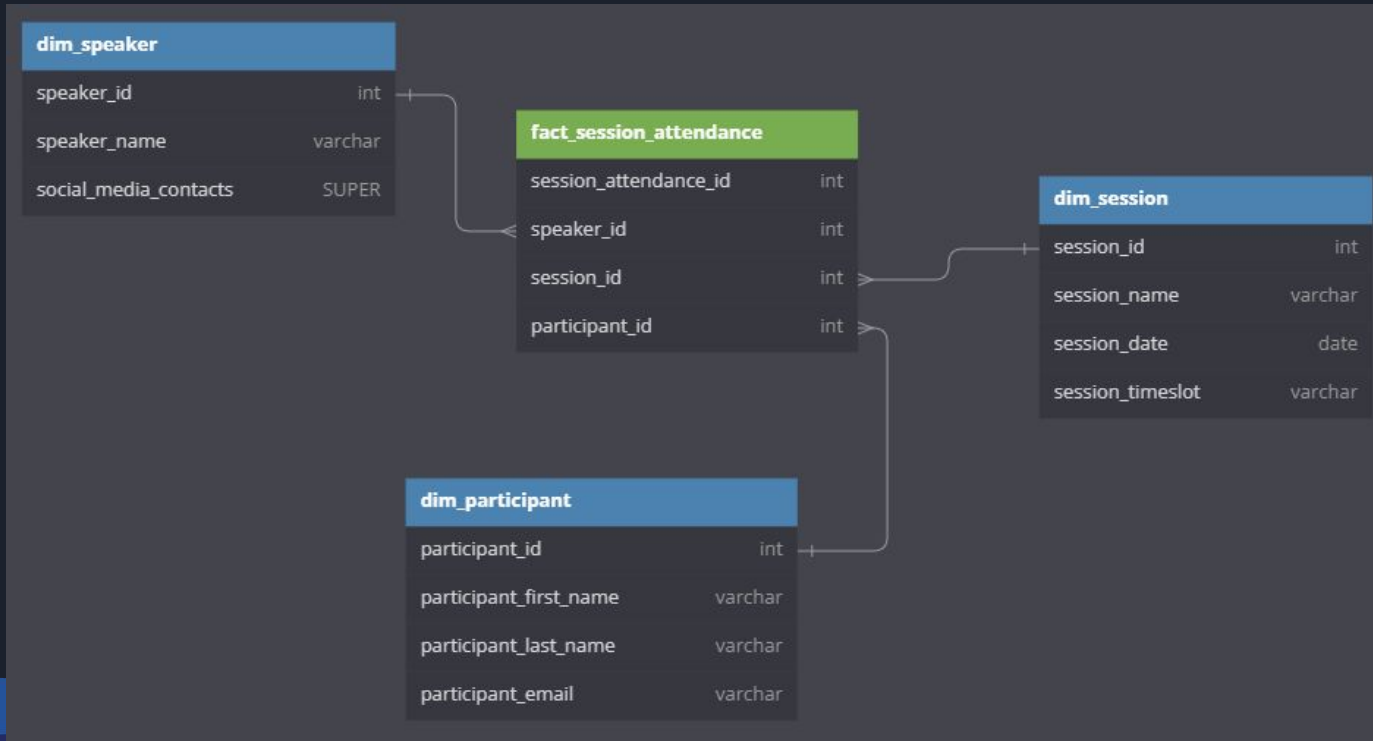
```
XN Hash Join DS_DIST_ALL_NONE (cost=0.38..0.74 rows=2 width=89)
  Hash Cond: ("outer".session_id = "inner".session_id)
    -> XN Hash Join DS_DIST_NONE (cost=0.25..0.57 rows=2 width=66)
      Hash Cond: ("outer".speaker_id = "inner".speaker_id)
        -> XN Hash Join DS_DIST_NONE (cost=0.13..0.40 rows=2 width=70)
          Hash Cond: ("outer".speaker_id = "inner".speaker_id)
            -> XN Seq Scan on dim_speaker_social dss (cost=0.00..0.11 rows=11 width=51)
            -> XN Hash (cost=0.12..0.12 rows=1 width=19)
              -> XN Seq Scan on dim_speaker ds (cost=0.00..0.12 rows=1 width=19)
                Filter: ((speaker_name)::text = 'John Smith'::text)
            -> XN Hash (cost=0.10..0.10 rows=10 width=8)
              -> XN Seq Scan on fact_session_attendance fsa (cost=0.00..0.10 rows=10 width=8)
            -> XN Hash (cost=0.10..0.10 rows=10 width=31)
              -> XN Seq Scan on dim_session d (cost=0.00..0.10 rows=10 width=31)
```



What if we want to
avoid a second table
altogether?



Option #2 - Our SUPER-hero



What does this look like in practice?

speaker_id	speaker_name	social_media_contacts
1	John Smith	[{"platform": "twitter", "username_url": "@johnsmith"}, {"platform": "linkedin", "username_url": "https://www.linkedin.com/in/johnsmith"}]
2	Jane Doe	[{"platform": "twitter", "username_url": "@janedoe"}, {"platform": "linkedin", "username_url": "https://www.linkedin.com/in/janedoe"}]
3	Michael Johnson	[{"platform": "twitter", "username_url": "@michaeljohnson"}, {"platform": "linkedin", "username_url": "https://www.linkedin.com/in/michaeljohnson"}]
4	Emily Brown	[{"platform": "twitter", "username_url": "@emilybrown"}, {"platform": "linkedin", "username_url": "https://www.linkedin.com/in/emilybrown"}]
5	David Davis	[{"platform": "twitter", "username_url": "@daviddavis"}, {"platform": "linkedin", "username_url": "https://www.linkedin.com/in/daviddavis"}]
6	Laura Miller	[{"platform": "twitter", "username_url": "@lauramiller"}, {"platform": "linkedin", "username_url": "https://www.linkedin.com/in/lauramiller"}]
7	Andrew Wilson	[{"platform": "twitter", "username_url": "@andrewwilson"}, {"platform": "linkedin", "username_url": "https://www.linkedin.com/in/andrewwilson"}]
8	Olivia Lee	[{"platform": "twitter", "username_url": "@olivialee"}, {"platform": "linkedin", "username_url": "https://www.linkedin.com/in/olivialee"}]
9	Daniel Clark	[{"platform": "twitter", "username_url": "@danielclark"}, {"platform": "linkedin", "username_url": "https://www.linkedin.com/in/danielclark"}]
10	Sophia Wang	[{"platform": "twitter", "username_url": "@sophiawang"}, {"platform": "linkedin", "username_url": "https://www.linkedin.com/in/sophiawang"}]





Let's pretty-print-ify this

```
[
  {
    "platform": "twitter",
    "username_url": "@johnsmith"
  },
  {
    "platform": "linkedin",
    "username_url": "https://www.linkedin.com/in/johnsmith"
  }
]
```



Our mission to alleviate joins is complete, but our users are going to HATE it.

Also, what about query performance?





Final Step - Materialized View

For ease of querying and consuming by other systems, a materialized view can be put over the data to simulate the data being stored in separate columns



Here's our goal

speaker_id	speaker_name	platform	username_url
1	John Smith	twitter	@johnsmith
1	John Smith	linkedin	https://www.linkedin.com/in/johnsmith
2	Jane Doe	twitter	@janedoe
2	Jane Doe	linkedin	https://www.linkedin.com/in/janedoe
3	Michael Johnson	twitter	@michaeljohnson
3	Michael Johnson	linkedin	https://www.linkedin.com/in/michaeljohnson
4	Emily Brown	twitter	@emilybrown
4	Emily Brown	linkedin	https://www.linkedin.com/in/emilybrown
5	David Davis	twitter	@davidddavis
5	David Davis	linkedin	https://www.linkedin.com/in/davidddavis
6	Laura Miller	twitter	@lauramiller
6	Laura Miller	linkedin	https://www.linkedin.com/in/lauramiller
7	Andrew Wilson	twitter	@andrewwilson
7	Andrew Wilson	linkedin	https://www.linkedin.com/in/andrewwilson
8	Olivia Lee	twitter	@olivialeee
8	Olivia Lee	linkedin	https://www.linkedin.com/in/olivialeee
9	Daniel Clark	twitter	@danielclark
9	Daniel Clark	linkedin	https://www.linkedin.com/in/danielclark
10	Sophia Wang	twitter	@sophiawang
10	Sophia Wang	linkedin	https://www.linkedin.com/in/sophiawang



Our materialized view code

```
CREATE MATERIALIZED VIEW vw_dim_speaker_social_contacts
  DISTSTYLE KEY
  DISTKEY (speaker_id)
  AS
SELECT speaker_id,
       speaker_name,
       REPLACE(socials.platform::varchar, '"', '')    platform,
       REPLACE(socials.username_url::varchar, '"', '') username_url
FROM conf_star_schema_v3.dim_speaker ds,
     ds.social_media_contacts socials;
```



What if we want the data in columns?

speaker_id ^	speaker_name ↕	twitter ↕	linkedin ↕	threads ↕
1	John Smith	@johnsmith	https://www.linkedin.com/in/johnsmith	<null>
2	Jane Doe	@janedoe	https://www.linkedin.com/in/janedoe	<null>
3	Michael Johnson	@michaeljohnson	https://www.linkedin.com/in/michaeljohnson	<null>
4	Emily Brown	@emilybrown	https://www.linkedin.com/in/emilybrown	<null>
5	David Davis	@daviddavis	https://www.linkedin.com/in/daviddavis	<null>
6	Laura Miller	@lauramiller	https://www.linkedin.com/in/lauramiller	<null>
7	Andrew Wilson	@andrewwilson	https://www.linkedin.com/in/andrewwilson	<null>
8	Olivia Lee	@olivialeee	https://www.linkedin.com/in/olivialeee	<null>
9	Daniel Clark	@danielclark	https://www.linkedin.com/in/danielclark	<null>
10	Sophia Wang	@sophiawang	https://www.linkedin.com/in/sophiawang	<null>



Our final materialized view code

```
CREATE MATERIALIZED VIEW vw_dim_speaker_social_contacts_pivoted
  DISTSTYLE KEY
  DISTKEY (speaker_id)
  AS
SELECT *
FROM (SELECT speaker_id,
              speaker_name,
              REPLACE(socials.platform::varchar, '"', ' '::varchar) platform,
              REPLACE(socials.username_url::varchar, '"', ' '::varchar) username_url
      FROM conf_star_schema_v3.dim_speaker ds,
           ds.social_media_contacts socials) PIVOT (
              max(username_url) FOR platform IN ('twitter', 'linkedin', 'threads')
      );
```



Final Takeaways

For distributed, massively parallel databases, JOIN statements can be very costly, especially if data is not distributed correctly

Using the correct distribution style and key will improve performance, but imposes limitations on how we query and design the rest of our model

To minimize joining entirely we can store sub-dimension data into JSON fields instead of their own tables

Materialized views can be put over the top to retain performance and increase usability



Thank you!

[@nickheidke](https://twitter.com/nickheidke) on Twitter?

Presentation materials available on: github.com/nickheidke



*Session
Feedback*



<https://that.land/3NOVbYe>





Resources

[Optimizing for Star Schemas on Redshift](#)

[Distribution and Sort Key Best Practices](#)

[PartiQL for Querying Semistructured Data](#)

[Using Materialized Views with the SUPER Datatype](#)

